

main.py - FastAPI Server & API Endpoints

Overview

File: server/app/main.py

Purpose: HTTP server that receives packets and provides web dashboards

Role: Central data receiver and API provider

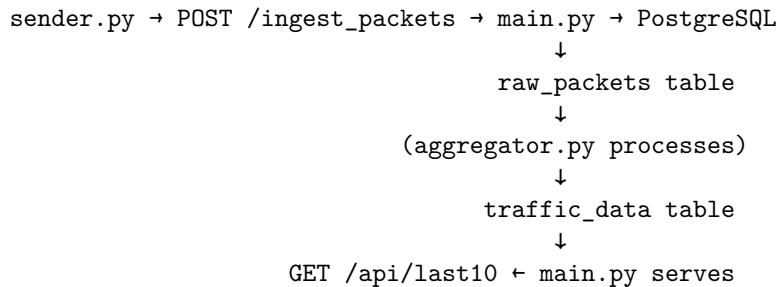
Runs as: Web server (via uvicorn)

What It Does

main.py is the FastAPI server that handles all HTTP communication:

1. **Receives** raw packet data from sender.py via POST /ingest_packets
 2. **Stores** packets in PostgreSQL database (raw_packets table)
 3. **Provides** web dashboards to view traffic flows
 4. **Serves** JSON APIs for programmatic access
 5. **Handles** ML predictions for aggregated flows (legacy endpoint)
-

Architecture



Key Components

1. Database Setup (Lines 20-82)

Tables Created:

raw_packets - Complete packet logs

```
raw_packets_table = Table(
    "raw_packets", metadata,
    Column("id", Integer, primary_key=True, autoincrement=True),
    Column("timestamp", Float),
```

```

        Column("interface", String),
        Column("src_ip", String),
        Column("dst_ip", String),
        Column("protocol", String),
        Column("length", Integer),
        Column("src_port", Integer, nullable=True),
        Column("dst_port", Integer, nullable=True),
        Column("tcp_flags", String, nullable=True),
        # ... 20+ more columns for DNS, HTTP, ICMP, etc.
        Column("inserted_at", DateTime, default=datetime.utcnow)
    )

```

Purpose: Audit trail and forensics - stores every packet captured

traffic_data - Aggregated flows with ML predictions

```

traffic_table = Table(
    "traffic_data", metadata,
    Column("id_num", Integer, primary_key=True, autoincrement=True),
    Column("dest_ip", String),
    Column("source_mac", String),
    Column("dest_mac", String),
    Column("packet_count", Integer),
    Column("packet_per_sec", Float),
    Column("byte_count", Integer),
    Column("byte_per_sec", Float),
    Column("tcp_flags", String),
    Column("connection_attempts", Integer),
    Column("unique_ports", Integer),
    Column("protocol", String),
    Column("predicted_label", String), # ML prediction
    Column("created_at", DateTime, default=datetime.utcnow)
)

```

Purpose: Stores aggregated network flows with ML intrusion detection labels

2. Pydantic Models (Lines 92-141)

Models define API request/response formats:

RawPacket - For /ingest_packets endpoint

```

class RawPacket(BaseModel):
    timestamp: float
    interface: str

```

```

src_ip: Optional[str] = None
dst_ip: Optional[str] = None
protocol: Optional[str] = None
length: int
# ... 25+ more optional fields

```

Validates incoming packet data from sender.py

Traffic - For /predict endpoint (legacy)

```

class Traffic(BaseModel):
    dest_ip: str
    source_mac: str
    dest_mac: str
    packet_count: int
    packet_per_sec: float
    # ... aggregated flow features

```

For manual ML prediction requests

3. ML Model Loading (Lines 85-90)

```

try:
    model = joblib.load("AI_model.pkl")
    logger.info("ML model loaded successfully")
except Exception as e:
    logger.error(f"Failed to load ML model: {e}")
model = None

```

Model file: server/models/AI_model.pkl

Used by: /predict endpoint for intrusion detection

Graceful handling: Server still runs if model missing

4. API Endpoints

POST /ingest_packets (Lines 172-191) Purpose: Receive batch of raw packets from sender.py

Request:

```
[
{
    "timestamp": 1701234567.123,
    "interface": "Wi-Fi",
    "packets": [
        {
            "src_ip": "192.168.1.100",
            "dst_ip": "192.168.1.101",
            "src_mac": "00:11:22:33:44:55",
            "dst_mac": "00:11:22:33:44:55",
            "protocol": "TCP",
            "length": 1460,
            "packet_count": 1000,
            "packet_per_sec": 100.0
        }
    ]
}
```

```

    "src_ip": "192.168.1.100",
    "dst_ip": "8.8.8.8",
    "protocol": "TCP",
    "length": 60,
    "src_port": 54321,
    "dst_port": 443,
    "tcp_flags": "SYN",
    "tcp_syn": true,
    ...
},
...
]

```

Response:

```
{
    "status": "success",
    "packets_received": 150,
    "message": "Packets stored in raw_packets table"
}
```

Code:

```

@app.post("/ingest_packets")
async def ingest_packets(packets: List[RawPacket]):
    with engine.begin() as conn:
        for packet in packets:
            ins = raw_packets_table.insert().values(**packet.dict())
            conn.execute(ins)

    logger.info(f"Inserted {len(packets)} raw packets")
    return {"status": "success", "packets_received": len(packets)}

```

POST /predict (Lines 193-214) Purpose: Predict traffic label for aggregated flow (legacy)

Request:

```
{
    "dest_ip": "8.8.8.8",
    "source_mac": "unknown",
    "dest_mac": "unknown",
    "packet_count": 150,
    "packet_per_sec": 5.0,
    "byte_count": 9000,
    "byte_per_sec": 300.0,
    "tcp_flags": "SYN,ACK",
}
```

```
        "connection_attempts": 3,  
        "unique_ports": 2,  
        "protocol": "TCP"  
    }  
}
```

Response:

```
{  
    "predicted_label": "Normal",  
    "data": { ... }  
}
```

Note: Mostly replaced by aggregator.py which automates this

GET / (Lines 219-267) **Purpose:** Web dashboard showing last 10 traffic flows

Returns: HTML page with auto-refreshing table

Features: - Shows last 10 flows from `traffic_data` table - Auto-refreshes every 5 seconds via JavaScript - Displays: ID, IPs, MACs, packet stats, predicted label

GET /alltraffic_page (Lines 269-317) **Purpose:** Web dashboard showing all traffic flows

Same as / but shows all flows (could be thousands)

GET /api/last10 (Lines 320-326) **Purpose:** JSON API for last 10 flows

Response:

```
[  
    {  
        "id_num": 123,  
        "dest_ip": "8.8.8.8",  
        "packet_count": 150,  
        "predicted_label": "Normal",  
        ...  
    },  
    ...  
]
```

Used by: Web dashboards, external tools

GET /api/alltraffic (Lines 328-334) Purpose: JSON API for all flows

Warning: Can return huge amounts of data

GET /api/raw_packets/last/{count} (Lines 336-342) Purpose: Get last N raw packets

Example: /api/raw_packets/last/100

Response:

```
[  
  {  
    "id": 5234,  
    "timestamp": 1701234567.123,  
    "src_ip": "192.168.1.100",  
    "dst_ip": "8.8.8.8",  
    ...  
  },  
  ...  
]
```

GET /health (Lines 344-351) Purpose: Health check for monitoring

Response:

```
{  
  "status": "healthy",  
  "model_loaded": true,  
  "timestamp": "2025-11-29T21:50:00"  
}
```

Used by: sender.py to verify server is up

Data Flow Example

Scenario: sender.py uploads 150 packets

1. sender.py sends HTTP POST:
POST /ingest_packets
Content-Type: application/json

[150 packet objects]
2. FastAPI receives request:

- Validates each packet against RawPacket model
 - Converts JSON → Python dictionaries
3. Database insertion:
- ```
with engine.begin() as conn:
 for packet in packets:
 INSERT INTO raw_packets VALUES (...)
```
4. Response sent:
- ```
{
    "status": "success",
    "packets_received": 150,
    "message": "Packets stored in raw_packets table"
}
```
5. aggregator.py (separate process):
- Queries new packets from raw_packets
 - Aggregates into flows
 - Runs ML predictions
 - Inserts into traffic_data
6. Web dashboard refreshes:
- Calls /api/last10
 - Gets aggregated flows with predictions
 - Updates HTML table
-

Configuration

Database Connection (Line 21):

```
DATABASE_URL = "postgresql://postgres:987456@localhost:5432/Traffic_Analyzer"
```

Change to:

```
DATABASE_URL = "postgresql://USER:PASSWORD@HOST:PORT/DATABASE"
```

Model Path (Line 86):

```
model = joblib.load("AI_model.pkl")
```

Make sure AI_model.pkl is in server/models/ directory

Running the Server

Start Server

```
cd server
uvicorn app.main:app --host 0.0.0.0 --port 8000
```

Options: - --host 0.0.0.0 - Listen on all network interfaces - --port 8000
- Port number - --reload - Auto-reload on code changes (development) -
--workers 4 - Multiple worker processes (production)

Production Deployment

```
uvicorn app.main:app \
    --host 0.0.0.0 \
    --port 8000 \
    --workers 4 \
    --log-level info \
    --access-log
```

Monitoring

Check Server Status

```
curl http://localhost:8000/health
```

Test Packet Ingestion

```
curl -X POST http://localhost:8000/ingest_packets \
    -H "Content-Type: application/json" \
    -d '['
        {
            "timestamp": 1701234567.123,
            "interface": "test",
            "src_ip": "192.168.1.1",
            "dst_ip": "8.8.8.8",
            "protocol": "TCP",
            "length": 60
        }
    ]'
```

View Web Dashboard

```
http://SERVER_IP:8000/
```

Query API

```
curl http://localhost:8000/api/last10
```

Troubleshooting

“Connection refused”

Cause: Server not running

Solution:

```
uvicorn app.main:app --host 0.0.0.0 --port 8000
```

“relation ‘raw_packets’ does not exist”

Cause: Database tables not created

Solution:

```
# Tables auto-created on startup:  
metadata.create_all(engine) # Line 81
```

Check database:

```
psql -U postgres -d Traffic_Analyzer  
\dt # List tables
```

“ML model not loaded”

Cause: AI_model.pkl missing

Solution: 1. Check file exists: ls server/models/AI_model.pkl 2. If missing, train model or copy from backup 3. Server still works without model (for ingestion only)

Slow Response Times

Causes: - Large database queries - Too many concurrent requests

Solutions: 1. Add database indexes: sql CREATE INDEX idx_timestamp
ON raw_packets(timestamp); CREATE INDEX idx_created_at ON
traffic_data(created_at DESC);

2. Use pagination:

```

# Instead of all:
limit(100).offset(0)

3. Scale with more workers:
uvicorn app.main:app --workers 8

```

Performance Tips

Database Connection Pooling

```

engine = create_engine(
    DATABASE_URL,
    pool_pre_ping=True,      # Already set
    pool_size=20,            # Add this
    max_overflow=40           # Add this
)

```

Batch Inserts (More Efficient)

```

# Current: Loop with individual inserts
for packet in packets:
    conn.execute(ins)

# Better: Bulk insert
conn.execute(
    raw_packets_table.insert(),
    [packet.dict() for packet in packets]
)

```

Add Response Compression

```

pip install fastapi[all]

from fastapi.middleware.gzip import GZipMiddleware
app.add_middleware(GZipMiddleware, minimum_size=1000)

```

Summary

main.py provides:

- HTTP API for receiving packet data
- Database storage for raw packets
- Web dashboards for viewing flows
- JSON APIs for programmatic access
- Health check endpoint
- ML prediction capability (with model)

It does NOT: Capture packets (that's sniffer.py)

Upload files (that's sender.py)

Aggregate flows (that's aggregator.py)

Dependencies: FastAPI, SQLAlchemy, PostgreSQL, ML model

Default Port: 8000

Web UI: http://SERVER_IP:8000/